

Generative AI MasterMind

A beginner Guide

1: FUNDAMENTALS OF GENERATIVE AI

1.1 What is Generative AI?

Definition: Generative AI is software that creates new content (text, images, code, audio and video) by predicting what should come next based on patterns it learned from millions of examples.

Analogies

It's like autocomplete on steroids. When your phone suggests the next word in a text message, that's a tiny version of what GenAI does — except GenAI can predict entire paragraphs, articles, or even code.

It's like a very well-read parrot. Imagine a parrot that has "read" the entire internet. It doesn't truly understand what it says, but it's incredibly good at producing text that *sounds* like it understands. It predicts word-by-word what a knowledgeable human would likely say next.

It's like a jazz musician improvising. A jazz player learns thousands of songs, then creates new music by combining patterns they've absorbed. GenAI does the same with text — mixing and remixing patterns to create something new.


Examples in Action

1. **ChatGPT writing an email** — You say "write a polite email declining a meeting," and it generates a full professional email because it's seen millions of similar emails.
2. **DALL-E creating images** — You describe "a cat wearing a spacesuit on Mars" and it creates an image by combining visual patterns of cats, spacesuits, and Mars.
3. **GitHub Copilot writing code** — You start typing a function name and it predicts the entire code block because it's seen similar patterns in millions of code repositories.
4. **ElevenLabs generating audio** — You start typing "record a podcast intro in a confident American accent saying..." and it predicts and generates the full natural voiceover because it's seen patterns in millions of hours of speech data.
5. **Runway ML generating video** — You start typing "animate a robot walking through a cyberpunk city" and it predicts and generates the full smooth video clip because it's seen patterns in millions of video frames and motion sequences.

The key thing to remember is... GenAI doesn't "think" or "know" things — it predicts what text/images should come next based on patterns. It's incredibly useful, but it's pattern-matching, not reasoning.


ANALOGIES

AUTOCOMPLETE ON STEROIDS




Predicts entire paragraphs, articles, code, not just words.

WELL-READ PARROT



Has "read" the internet, mimics understanding by predicting what a human would say.

JAZZ MUSICIAN IMPROVISING



Learns patterns and remixes them to create something new.

EXAMPLES IN ACTION

CHATGPT WRITING AN EMAIL

Generates professional email based on millions of similar examples.

DALL-E CREATING IMAGES

Creates image by combining visual patterns (cats, spacesuits, Mars).

GITHUB COPILOT WRITING CODE

Predicts entire code blocks based on millions of repositories.

ELEVENLABS GENERATING AUDIO

Generates natural voiceover from speech data patterns.

RUNWAY ML GENERATING VIDEO

Generates video from motion sequence patterns.

RUNWAY ML GENERATING VIDEO

Generates video from region motion sequence patterns.



KEY TAKEAWAY

GenAI doesn't 'think' or 'know' — it predicts based on patterns. It's incredibly useful, but it's pattern-matching, not reasoning.

1.2 Regular AI vs. Generative AI

One-sentence difference: Regular AI finds, sorts, or decides — Generative AI creates something new.

The Core Difference

Regular AI	Generative AI
Finds existing answers	Creates new content
Chooses from what exists	Makes something that never existed
"Here's what I found"	"Here's what I made"

Analogy: Library vs. Author


Regular AI is like a librarian. You ask a question, they search the shelves, and hand you a book that already exists. They're great at finding, organizing, and recommending — but they don't write new books.

Generative AI is like an author. You give them a topic, and they write a brand new book that never existed before. It might be inspired by books they've read, but the output is original.

Examples Side-by-Side

Task	Regular AI Does This	Generative AI Does This
Search	Google finds existing webpages	ChatGPT writes a new answer
Images	Google Images finds existing photos	DALL-E creates a new image
Music	Spotify recommends existing songs	Suno creates a new song
Shopping	Amazon recommends products you might like	AI designs a new product
Email	Spam filter sorts into spam/not spam	AI writes a new email for you
Driving	Tesla decides: brake or accelerate	— (not a generation task)

Simple Diagram



Types of Regular AI (Non-Generative)

Type	What It Does	Example
Classification	Sorts things into categories	Email → Spam or Not Spam
Recommendation	Suggests existing items	"You might also like..."
Prediction	Forecasts outcomes	"80% chance of rain tomorrow"
Detection	Finds patterns/anomalies	Fraud detection on credit cards
Recognition	Identifies things	Face unlock on your phone

None of these create — they all analyze, sort, or choose.

Why This Matters

Regular AI: Great for automation, decisions, and finding things. Been around for decades.

Generative AI: New capability (mainstream since ~2022). Creates text, images, code, audio, video. This is the "ChatGPT revolution."

The key thing to remember is...

Regular AI = **Picker** (chooses from what exists)


Generative AI = **Maker** (creates something new)

Both are AI. Generative AI is just a specific type that focuses on creation rather than selection.

1.3 Inside the Black Box: How Generative AI Works (Concise)

One-liner: GenAI predicts the most likely next word, over and over, based on patterns learned from billions of examples.

The 5 Stages



Stage 1: Training

What: AI reads billions of webpages/books and learns patterns, not facts.

Analogy: A student who reads every library book but memorizes how sentences flow, not specific facts.

```
Sees: "Capital of France is Paris"  
      "Capital of Japan is Tokyo"  
Learns: "Capital of [country] is [city]" pattern
```

Stage 2: Parameters

What: Patterns stored as billions of numbers (weights).

Analogy: Like muscle memory — you can't explain how you ride a bike, but your body "knows."

Model	Parameters
GPT-2	1.5B
GPT-4	~1.8T

Stage 3: Tokens

What: Text broken into chunks (usually word pieces).

Analogy: Lego bricks — smaller pieces AI can mix and match.

```
"Explain photosynthesis" →  
["Explain", " photo", "synth", "esis"]
```

Why it matters: Context window, API cost, and speed all measured in tokens.

Stage 4: Attention

What: AI decides which tokens matter most for the answer.

Analogy: Highlighting key words in a sentence.

```
"The Eiffel Tower, built in 1889, is in which city?"  
  ▲ HIGH ATTENTION                ▲ HIGH ATTENTION
```

This is the "T" in GPT — Transformer — the breakthrough that lets AI see all words at once.


Stage 5: Prediction

What: Predict next token → add it → repeat thousands of times.

Analogy: Super-autocomplete.

```
"Best language for beginners is" →  
Python (42%) | JavaScript (18%) | Scratch (12%)...  
→ Picks "Python" → Repeats for next word
```

Temperature (Creativity Dial)




Common Myths

Myth	Reality
"AI understands"	It predicts likely responses
"AI has a fact database"	Knowledge = compressed patterns
"AI remembers our chat"	Re-reads entire history each time

Putting It All Together: A Complete Example


Here's what actually happens when you send a prompt to GenAI:



1.4 Types of AI Models: Foundational, Proprietary, Open Source

One-liner: Foundational models are the "engines" — proprietary means the company keeps the engine secret, open source means anyone can see and modify it. Products like Perplexity and Gamma are "cars" built using these engines.

The Three Categories



1. Foundational Models

What: Large pre-trained models that serve as the base "intelligence" — trained on massive data, general-purpose.

Analogy: The engine in a car. You don't build an engine from scratch to make a taxi service — you buy one and

build around it.

Characteristic	Description
Training cost	\$10M - \$100M+
Data	Trillions of tokens
Purpose	General-purpose base to build on
Who builds	Big tech companies

2. Proprietary Models (Closed Source)

What: Foundational models where the weights, training data, and code are kept secret. Access via API only.

Analogy: Coca-Cola's secret recipe — you can buy the drink, but you can't see how it's made.

Model	Company	Access
GPT-4, GPT-4o	OpenAI	API only
Claude 3.5/4	Anthropic	API only
Gemini	Google	API only

Pros: Often highest performance, managed infrastructure

Cons: No customization, vendor lock-in, data privacy concerns

3. Open Source Models

What: Foundational models where weights (and sometimes code/data) are publicly available. Download and run yourself.

Analogy: A recipe posted online — you can make it at home, modify it, sell your version.

Model	Company	Parameters
Llama 3.1	Meta	8B - 405B
Mistral	Mistral AI	7B - 8x22B
Falcon	TII	7B - 180B
Gemma	Google	2B - 27B

Pros: Free, customizable, data stays private, no API costs

Cons: Need infrastructure, often slightly lower performance

Where Products Fit: The Application Layer

These companies don't build foundational models — they build products ON TOP of them:

APPLICATION LAYER

Product	What It Does	Models Used
 Perplexity	AI search engine	• GPT-4 • Claude • own fine-tuned
 Gamma	AI presentations	• GPT-4 • Claude
 Jasper	Marketing copy	• GPT-4 • Claude • own fine-tuned
 Notion AI	Workspace AI	• GPT-4 • Claude
 GitHub Copilot	Code assistant	GPT-4 + Codex

Analogy:


Layer	Car Analogy	AI Example
Foundational Model	Engine (Toyota, Ford)	GPT-4, Llama, Claude
Application	Car service (Uber, Lyft)	Perplexity, Gamma
End User	Passenger	You

Uber doesn't manufacture engines — they build a service using cars with existing engines. Same with Perplexity using GPT-4.

Quick Comparison

Aspect	Proprietary	Open Source	App Layer
Example	GPT-4, Claude	Llama, Mistral	Perplexity, Gamma
Cost to use	Pay per API call	Free (but need hardware)	Subscription
Customization	Limited	Full control	Use as-is
Data privacy	Sent to provider	Stays with you	Sent to provider
Best for	Quick start, best performance	Privacy, cost control, customization	Non-technical users

Visual: Who Builds What



1.5 The Three Critical Limitations

GenAI has three major weaknesses: it can make things up (hallucinations), its knowledge has a cutoff date, and it can only "remember" a limited conversation.

Limitation 1: Hallucinations

It's like a confident friend who never admits they don't know something. Instead of saying "I don't know," GenAI will confidently make up an answer that sounds completely plausible but is totally wrong.

Examples:

- Citing research papers that don't exist
- Inventing historical facts with specific (fake) dates
- Creating fake statistics that sound believable

Limitation 2: Knowledge Cutoff

It's like talking to someone who's been in a coma since a certain date. GenAI was trained on data up to a specific point. It has no idea what happened after that — no recent news, no new discoveries, no current events.

Example: If the cutoff is January 2024, asking "Who won the 2024 Super Bowl?" will get you a guess, not a fact.

Limitation 3: Context Window (Memory Limits)

It's like talking to someone with short-term memory loss. GenAI can only "remember" a certain amount of your conversation. Share a 100-page document? It might forget the beginning by the time it reaches the end.

Conversation Start ←-----→ Context Limit
[Remembers this clearly] [Starts forgetting]

The key thing to remember is... Always verify important facts, check if information might be outdated, and break long documents into smaller chunks.

1.6 VERIFY Framework (Using AI Responsibly)

One-sentence definition: VERIFY is a 6-step checklist to make sure you're using AI safely and getting accurate results.

The Framework

Letter	Meaning	It's Like...
V	Validate outputs against trusted sources	Double-checking Wikipedia with textbooks
E	Exclude sensitive data from prompts	Not sharing passwords on social media
R	Review for bias and accuracy	Getting a second opinion from a friend
I	Iterate prompts to improve results	Asking follow-up questions until you understand
F	Flag uncertainty in AI responses	Putting a "?" next to anything you're not sure about
Y	Your judgment is final	You're the driver, AI is just GPS

Analogy

VERIFY is like being a newspaper fact-checker. Journalists don't publish the first thing a source tells them. They verify with multiple sources, remove sensitive info, check for bias, ask follow-up questions, note uncertainties, and ultimately make the final call on what to publish.

The key thing to remember is... AI is a powerful assistant, but you're responsible for the final output. Always verify before trusting.

2: PROMPT ENGINEERING

2.1 Meta & COSTAR Prompting Frameworks

One-sentence definition: Instead of writing prompts yourself, use advanced frameworks like meta-prompting (asking AI to write the prompt) or COSTAR (structured 6-part prompts) for world-class results.

Meta Prompting

What it is: A higher-level technique where you ask the AI to write the prompt for you instead of writing the final instruction yourself.

How it works:

"I want write to [Product Requirements Documents (PRD) for an account payable automation solution].

Act as an expert Prompt Engineer. Please write the best possible prompt that I can use to get a world-class response from you.

Include any necessary context, role-playing personas, or constraints that will make the output high-quality."

When to use: When you're unsure how to structure a complex request, or when you want the AI to help you craft the perfect prompt before executing the task.

COSTAR Framework

What it is: COSTAR is a prompt engineering framework used to structure instructions for Large Language Models (LLMs) like GPT-4 or Claude. Instead of asking a vague question, you break your request into six specific components to ensure the AI understands exactly what you need.

It stands for:

- **C** - Context (Background information)
- **O** - Objective (The specific task)
- **S** - Style (Writing style or persona)
- **T** - Tone (Emotional attitude)
- **A** - Audience (Who is reading this?)
- **R** - Response (Format of the output)

Example: Writing a Rejection Email

Without COSTAR (Bad):

```
"Write a rejection email to a candidate."
```

With COSTAR (Good):

Context: I am a hiring manager at a tech startup. We interviewed a junior developer named Alex who was great culturally but lacked the necessary Python experience.

Objective: Write a rejection email that encourages them to apply again in 6 months after learning more Python.

Style: Professional but human, like a mentor speaking to a student.

Tone: Empathetic and encouraging, not cold or corporate.

Audience: A young college graduate who is eager to learn.

Response: Format as a plain text email body, ready to copy-paste.

Why the COSTAR version wins: The AI knows why the candidate was rejected (Context), specific advice to give (Objective), and that it shouldn't sound like a robot (Tone/Style).

When to Use Meta-Prompt vs COSTAR Framework

Here's how to decide which strategy to use:

Scenario A: I need a specific deliverable (e.g., A JSON object, a SQL query, or a cold email)

→ **Use COSTAR** — You know what you want, you just need to structure the request clearly.

Scenario B: I'm not sure how to structure my request or ask

→ **Use Meta-Prompting** — Let the AI help you design the perfect prompt first, then execute it.

2.2 Most Commonly Used Prompt Engineering Techniques

This section provides a comprehensive exploration of 12 prompt engineering techniques with practical examples, use cases, and when to apply each one.

1. Zero-Shot Prompting

When to Use:

- Simple, well-defined tasks
- LLM has strong prior knowledge of the domain

- No examples available yet
- Quick prototyping or testing

When NOT to Use:

- Complex formatting requirements
- Domain-specific jargon or style
- Tasks requiring specific output structure
- Ambiguous instructions where examples would clarify

Structure:

```
[Clear instruction] + [Task details] + [Output format]
```

Bad Example:

```
Write something about Python error handling.
```

Problem: Vague scope, no format, no specific angle

Good Example:

```
Explain Python's try-except-finally block in 3 sentences.  
Include one code snippet showing proper exception handling.  
Target audience: intermediate developers.
```

Why: Specific scope, format, length, and audience defined

2. Few-Shot Prompting

When to Use:

- Specific formatting needed (JSON, tables, templates)
- Style matching required
- Classification tasks
- Pattern recognition needed
- Domain-specific transformations

When NOT to Use:

- LLM already understands the format
- Examples are inconsistent or conflicting
- Task is self-explanatory
- Limited context window (examples consume tokens)

Structure:

```
[Task description]
```

```
Example 1:
```

```
Input: [example input]
```

```
Output: [example output]
```

```
Example 2:
```

```
Input: [example input]
```

```
Output: [example output]
```

```
Now apply to:
```

```
Input: [actual input]
Output:
```

Bad Example:

```
Convert to JSON:
Name: John, Age: 30
Name: Sarah, Age: 25

Convert: Name: Mike, Age: 40
```

Problem: Inconsistent formatting, no field definitions

Good Example:

```
Extract structured data from user messages:

Example 1:
Input: "I need a flight to NYC on Dec 15th"
Output: {"type": "flight", "destination": "NYC", "date": "2024-12-15"}

Example 2:
Input: "Book hotel in Paris for 3 nights starting Jan 5"
Output: {"type": "hotel", "location": "Paris", "duration": "3 nights", "start_date": "2025-01-05"}

Extract from: "Reserve rental car in Miami from Feb 10-12"
Output:
```

Why: Consistent structure, clear field mapping, diverse examples

3. Chain-of-Thought (CoT)

When to Use:

- Multi-step reasoning required
- Mathematical or logical problems
- Complex decision-making
- Debugging or analysis tasks
- Need to verify reasoning path

When NOT to Use:

- Simple retrieval or lookup tasks
- When only final answer matters (tokens are precious)
- Real-time/low-latency requirements
- Creative writing (can make it mechanical)

Structure:

```
[Task] + "Think step-by-step" OR "Show your reasoning before answering"
```

Bad Example:

```
Calculate the ROI of our marketing campaign. Think step-by-step.
```

Problem: Missing data, unclear what steps to show

Good Example:

Calculate marketing ROI with this data:

- Campaign cost: \$10,000
- Revenue generated: \$45,000
- Attribution: 60% direct, 40% assisted

Think step-by-step:

1. Calculate attributed revenue
2. Subtract campaign cost
3. Compute ROI percentage
4. Explain if this meets our 3x ROI target

Why: Data provided, specific steps outlined, clear success criteria

4. ReAct (Reasoning + Acting)

When to Use:

- Multi-tool workflows (search → analyze → generate)
- Agent-based systems
- Tasks requiring external information
- Dynamic decision-making (if X then Y)
- N8N workflows with branching logic

When NOT to Use:

- Single-step tasks
- No external tools available
- Linear workflows (use prompt chaining instead)
- When reasoning overhead isn't needed

Structure:

```
Task: [goal]
Available tools: [list]

For each step:
Thought: [reasoning]
Action: [tool_name(parameters)]
Observation: [result]
... (repeat until)
Answer: [final response]
```

Bad Example:

```
Search for Python tutorials and summarize them.
Tools: web_search
```

Problem: No reasoning structure, unclear iteration

Good Example:

```
Find the latest Python 3.12 async features and create a code example.

Available tools: web_search(query), web_fetch(url)

Thought: I need recent documentation about Python 3.12 async features
Action: web_search("Python 3.12 async new features")
Observation: [results show TaskGroups and asyncio improvements]

Thought: Need official documentation for accuracy
Action: web_fetch("https://docs.python.org/3.12/whatsnew/3.12.html")
```


Observation: [full docs retrieved]

Thought: Now I can create accurate example

Answer: [code example with TaskGroups]

Why: Clear thought → action → observation loop, explicit reasoning

5. Role-Based Prompting

When to Use:

- Domain expertise needed (legal, medical, technical)
- Specific perspective required
- Tone/style customization
- Technical documentation
- Code review or architecture decisions

When NOT to Use:

- Generic tasks where role doesn't add value
- Over-specification that constrains creativity
- When expertise might introduce bias

Structure:

You are a [specific role with expertise].
Your task is [goal].
Approach this as [persona] would, focusing on [domain aspects].

Bad Example:

You are an expert. Write Python code for data processing.

Problem: Vague expertise, no specific domain knowledge leveraged

Good Example:

You are a senior data engineer specializing in real-time ETL pipelines using Python, Apache Kafka, and AWS.

Review this data ingestion code for:

- Scalability issues (target: 10K events/sec)
- Error handling in streaming contexts
- Memory leaks in long-running processes
- Best practices for backpressure handling

[code here]

Why: Specific expertise, measurable criteria, domain context

6. Prompt Chaining

When to Use:

- Complex workflows (RAG pipelines)
- Quality > speed
- Each step needs different specialization
- Intermediate outputs need validation
- N8N automation sequences

When NOT to Use:

- Simple single-step tasks
- Real-time requirements
- Limited API calls budget
- When outputs don't build on each other

Structure:

```
Prompt 1: [specialized task A] → Output A
Prompt 2: Use Output A to [specialized task B] → Output B
Prompt 3: Use Output B to [final task] → Final Output
```

Bad Example:

```
Chain:
1. Summarize this article
2. Make it better
3. Add hashtags
```

Problem: Vague objectives, no clear specialization per step

Good Example:

```
CHAIN 1 - Query Rewriting:
User query: "python async stuff"
Rewrite as 3 specific search queries optimized for technical documentation.
Output: {{search_queries}}

CHAIN 2 - Retrieval Ranking:
Retrieved docs: {{retrieved_chunks}}
Rank by relevance to: {{search_queries}}
Return top 5 with confidence scores.
Output: {{ranked_docs}}

CHAIN 3 - Synthesis:
Context: {{ranked_docs}}
Original query: "python async stuff"
Generate answer with:
- Code examples
- Citations (doc IDs)
- Confidence score (0-1)
Output format: JSON
```

Why: Each step has clear input/output, specialization, structured flow

7. Constraint-Based Prompting

When to Use:

- Need to prevent specific behaviors
- Format compliance (legal, regulatory)
- Brand guidelines enforcement
- API response requirements
- Content moderation

When NOT to Use:

- Over-constraining kills creativity
- Constraints conflict with each other

- Simple tasks where flexibility is fine

Structure:

[Task]

MUST include: [requirements]
MUST NOT include: [prohibitions]
Format: [structure]
Length: [bounds]

Bad Example:

Write a blog post. Keep it professional.

Problem: "Professional" is subjective, no measurable constraints

Good Example:

Write a product description for enterprise RAG platform.

MUST include:
- One quantified benefit (e.g., "40% faster retrieval")
- One integration mention (Pinecone, Weaviate, or ChromaDB)
- Use case example from finance OR healthcare

MUST NOT include:
- Marketing fluff ("revolutionary", "game-changing")
- Unverified claims
- Technical jargon not explained
- More than 150 words

Format: 3 paragraphs with subheadings
Tone: Technical but accessible

Why: Specific inclusions/exclusions, measurable limits, clear structure

8. Self-Consistency Sampling

When to Use:

- High-stakes decisions (medical, legal, financial)
- Ambiguous problems with multiple valid approaches
- Need confidence estimation
- Quality > cost

When NOT to Use:

- Deterministic tasks (formatting, data extraction)
- Budget constraints (requires multiple API calls)
- Real-time systems
- Tasks with single correct answer

Structure:

Generate N different reasoning paths for:
[problem]

Compare answers and:
1. Identify consensus answer
2. Note divergent reasoning

3. Provide confidence score based on agreement

Bad Example:

Give me 3 answers to: What's the capital of France?

Problem: Deterministic question, no value in multiple samples

Good Example:

A startup has \$500K funding and two options:
Option A: Hire 5 engineers (\$400K/year) + minimal marketing
Option B: Hire 2 engineers (\$160K/year) + \$200K marketing + freelancers

Generate 3 independent analyses considering:

- 18-month runway
- B2B SaaS product (0 revenue currently)
- Technical complexity requires strong team

For each analysis:

1. Compute runway for each option
2. Assess risk factors
3. Recommend option

Then: Compare recommendations and explain consensus or disagreement.

Why: Ambiguous problem, multiple valid approaches, genuine uncertainty

9. Meta-Prompting (Prompt Generation)

When to Use:

- Building prompt templates for non-technical users
- Creating N8N workflow prompt libraries
- A/B testing prompt variations
- Teaching prompt engineering

When NOT to Use:

- Direct task execution (just do the task instead)
- Over-engineering simple prompts
- When you already have proven prompts

Structure:

Generate a prompt for [use case] that:

- Handles [input type]
- Produces [output format]
- Includes [specific constraints]
- Uses [technique name] technique

Provide 3 variations: basic, intermediate, advanced.

Bad Example:

Make me a prompt for summarizing.

Problem: No specifications, no use case context

Good Example:

Generate a prompt template for N8N workflow that:

Use case: Summarize customer support tickets

Input: {{ticket_text}} (variable length, 50-500 words)

Output: JSON with keys: {summary, sentiment, priority, suggested_action}

Requirements:

- Use few-shot learning with 2 examples
- Include constraint: summary max 50 words
- Add instruction for extracting action items
- Specify priority scale: low/medium/high/urgent

Provide:

1. The complete prompt template
2. Example with sample ticket
3. Expected output for the example

Why: Complete specifications, structured deliverable, actionable template

10. Negative Prompting

When to Use:

- Preventing known unwanted patterns
- Removing verbose AI behaviors (apologies, disclaimers)
- Content filtering
- Style refinement

When NOT to Use:

- Over-constraining creative tasks
- When positive instructions are clearer
- Adding too many negatives (confusing)

Structure:

[Task]

DO NOT:

- [behavior 1]
- [behavior 2]
- [behavior 3]

Bad Example:

Write code. Don't make mistakes.

Problem: "Mistakes" is too vague

Good Example:

Generate Python function for API rate limiting.

DO NOT:

- Include apologies or explanatory preamble
- Use time.sleep() (use asyncio instead)
- Add print statements (use logging module)
- Write comments explaining obvious code
- Include error handling for network issues (caller handles this)

DO:

- Use decorator pattern

- Include type hints
- Return remaining quota in response

Why: Specific behaviors prevented, balanced with positive instructions

11. Output Structuring

When to Use:

- API integrations (N8N, Zapier)
- Data parsing requirements
- Multi-field extraction
- Downstream processing in pipelines

When NOT to Use:

- Human-only consumption (prose is fine)
- Format complexity exceeds task value
- LLM struggles with strict JSON (use XML instead)

Structure:

[Task]

Return ONLY valid JSON with this exact structure:

```
{
  "field1": "type",
  "field2": ["array", "of", "items"],
  "nested": {
    "field3": "value"
  }
}
```

No markdown code blocks, no explanations, pure JSON only.

Bad Example:

Extract info from this email and give me JSON:
[email content]

Problem: No schema defined, "info" is ambiguous

Good Example:

Extract meeting details from email:

[email: "Let's meet Tuesday at 2pm at Coffee Shop on Main St to discuss Q4 budget"]

Return ONLY this JSON structure (no markdown, no explanations):

```
{
  "meeting_date": "YYYY-MM-DD or null if not specific",
  "meeting_time": "HH:MM or null",
  "location": "string or null",
  "attendees": ["array of names mentioned"],
  "topics": ["array of discussion topics"],
  "action_items": ["array or empty"],
  "confidence": 0.0-1.0
}
```

If any field cannot be determined, use null. Always include confidence score.

Why: Exact schema, data types, null handling, parsing rules clear

12. Iterative Refinement Prompting

When to Use:

- Complex deliverables (reports, documentation)
- Collaborative AI workflows
- Learning user preferences
- Quality improvement cycles

When NOT to Use:

- Simple one-shot tasks
- Stateless API calls (no conversation memory)
- Time-sensitive requests

Structure:

```
Iteration 1: Generate [draft]
Iteration 2: Critique your output for [criteria]
Iteration 3: Revise based on critique
```

Bad Example:

```
Write an article. Then make it better.
```

Problem: "Better" undefined, no critique criteria

Good Example:

```
Task: Create RAG system architecture diagram description

Iteration 1: Write technical description of components
Iteration 2: Self-critique against these criteria:
- Are all data flows explained?
- Is latency at each stage mentioned?
- Are error handling paths covered?
- Would a senior engineer spot gaps?

Iteration 3: Revise description addressing each critique point

Iteration 4: Add one-paragraph "Scaling Considerations" section
```

Why: Specific improvement criteria, staged refinement, measurable progress

Combining Techniques (Advanced)

RAG Pipeline Example:

1. Few-Shot + Output Structuring: Query rewriting
2. Constraint-Based: Filter retrieval results
3. CoT: Rank documents by relevance
4. Role-Based + Negative: Generate answer (expert tone, no disclaimers)
5. Output Structuring: Format as JSON with citations

This combination creates a production-ready RAG response in a structured, repeatable way.

Practice Exercise

Build a RAG pipeline prompt that:

1. Rewrites user query for better retrieval
2. Ranks retrieved chunks by relevance
3. Synthesizes answer with citations
4. Returns JSON with confidence score

Solution Framework:

```
# Step 1: Query Rewriting (Few-Shot + Output Structuring)
Original query: {{user_query}}

Rewrite into 3 optimized search queries:
Example: "python async" → ["python asyncio library", "async await python syntax", "python asynchronous programming tutorial"]

Output: ["query1", "query2", "query3"]

# Step 2: Ranking (CoT + Role-Based)
You are a search relevance engineer.

Retrieved chunks: {{chunks}}
Queries: {{rewritten_queries}}

Think step-by-step:
1. Score each chunk against each query (0-1)
2. Calculate weighted average
3. Rank top 5 chunks

Output JSON: [{"chunk_id": "x", "score": 0.95, "reason": "..."}, ...]

# Step 3: Synthesis (Constraint-Based + Negative)
Context: {{top_chunks}}
Original query: {{user_query}}

Generate answer:
MUST include: Code examples, citations [1], [2]
MUST NOT: Apologize, use disclaimers, exceed 200 words

# Step 4: Final Output (Output Structuring)
Return ONLY this JSON:
{
  "answer": "string",
  "citations": [{"id": 1, "chunk_id": "x", "relevance": 0.95}],
  "confidence": 0.85,
  "suggested_followups": ["question1", "question2"]
}
```

Key Takeaways

1. **Match technique to task complexity** - Don't over-engineer simple prompts
2. **Token costs matter** - CoT and chaining are expensive; use strategically
3. **Examples beat instructions** - Few-shot works when explanations fail
4. **Structure enables automation** - JSON output is essential for N8N workflows
5. **Combine techniques** - Real systems need 3-5 techniques working together
6. **Test iteratively** - Start simple, add complexity based on failures
7. **Document patterns** - Save successful prompts as templates

Additional Resources

- **Anthropic Prompt Engineering Guide:** <https://docs.claude.com/en/docs/build-with-claude/prompt-engineering/overview>
- **N8N AI Documentation:** For workflow integration patterns

- **LangChain Prompts:** Template library for inspiration

Prompt Engineering Techniques Guide - Version 1.0 - January 2026

3: CONTEXT ENGINEERING

[Content to be added - This section will cover RAG systems, context windows, chunking strategies, embeddings, vector databases, and context optimization techniques]

4: MODEL CONTEXT PROTOCOL (MCP)


[Content to be added - This section will cover MCP fundamentals, implementation patterns, use cases, and integration strategies]

5: THE FIVE AI ARCHITECTURES

5.1 The Five AI Architectures

There are 5 different ways to build AI solutions, ranging from simple chat to fully autonomous AI teams — choosing the right one depends on your task complexity.


The Spectrum



Architecture 1: Basic LLM Chat

Direct conversation: with AI using only the knowledge it was trained on — like texting a very smart friend who has read a lot of books.

It's like: A walking encyclopedia you can have a conversation with. Great for general knowledge, but it hasn't read your company's documents or today's news.



Best for: Brainstorming, drafting content, explaining concepts, quick questions where general knowledge is enough.

Limitations: Doesn't know your specific data, can hallucinate facts, knowledge cutoff.


Examples:

- "Help me brainstorm product names"
- "Explain machine learning to a 10-year-old"
- "Draft a thank-you email"

Architecture 2: RAG (Retrieval-Augmented Generation)

AI that first searches your documents: to find relevant information, then generates answers based on what it found — like an assistant who checks your files before answering.

It's like: A librarian who searches the right books first, then answers your question using what they found. They quote real sources instead of guessing.



Best for: Q&A over company documents, customer support with accurate answers, research across large document collections.

Key components explained simply:


Component	What It Does	Analogy
Document Loader	Reads your files	Scanner at a library
Chunker	Splits docs into searchable pieces	Cutting a book into chapters
Embeddings	Converts text to searchable format	Creating an index
Vector Database	Stores and finds chunks	The library's card catalog
LLM	Writes the final answer	The librarian explaining what they found

The key thing to remember is... RAG = AI that checks YOUR documents before answering, reducing hallucinations and giving sourced answers.

Architecture 3: AI Workflow (Orchestrated Pipelines)

Multiple AI steps chained together: in a fixed sequence, where each step does one specific job — like a factory assembly line where AI handles certain stations.

It's like: A car assembly line. Step 1 adds the frame, Step 2 adds the engine, Step 3 paints it. Each step has one job, in a specific order, every time.



Best for: Repeatable processes, document processing, content pipelines where you know exactly what steps are needed.

Key characteristics:

Aspect	Description
Flow	Fixed path — same every time
Control	You design the steps
Reliability	High — predictable results
Flexibility	Low — can't adapt to surprises

Example workflows:

Weekly competitor monitor:

Timer (weekly) → Fetch competitor website → Extract changes (AI) → Compare to last week (AI) → Format report → Send to Slack

Cold Outreach:

Timer (daily) → Get company names(AI)
→ Send Email Outreach emails (AI)

Tools: N8N, Make, Zapier, Power Automate


The key thing to remember is... Workflows are for *predictable* multi-step processes where you know exactly what needs to happen in what order.

Architecture 4: AI Agent

AI that uses tools and decides its own next steps to complete a task — like a expert software engineer who figures out what to do instead of being told every step.

It's like: Giving someone a goal ("research and book me a flight to Tokyo") instead of step-by-step instructions.

They decide: Should I check prices first? Search multiple airlines? What dates work best? They use tools (websites, calendars) and adapt as they go.



Best for: Research tasks, problems where you don't know the exact steps in advance, tasks requiring multiple tool uses.

Key characteristics:

Aspect	Description
Flow	Dynamic — AI chooses next step
Control	Goal-directed (you give the goal)
Reliability	Medium — can get stuck or loop
Flexibility	High — adapts to what it finds

Workflow vs. Agent:

Workflow	Agent
"Do step 1, then 2, then 3"	"Achieve this goal"
You plan the steps	AI plans the steps
Same path every time	Different path based on situation
Like a recipe	Like a personal assistant


The key thing to remember is... Agents are for tasks where you know the *goal* but not the exact steps. You give them tools and let them figure it out.

Architecture 5: Agentic AI (Multi-Agent Systems)

Multiple specialized AI agents working together as a team, each with different expertise, collaborating to tackle complex projects — like an AI company with different departments.

It's like: A film production crew. You have a director (orchestrator), writer (content agent), cinematographer (visual

agent), and editor (review agent). Each is an expert in their role, they pass work to each other, and together they create something none could alone.



Common team patterns:

Pattern	How It Works	Example
Pipeline	Agents work in sequence	Research → Write → Edit → Publish
Supervisor	One agent assigns tasks	Manager agent delegates to specialists
Debate	Agents argue different sides	Pro agent vs. Con agent for decisions
Swarm	Agents work in parallel	5 researchers each tackle different sources

Best for: Complex projects requiring multiple expertise areas, tasks benefiting from review/critique cycles, situations where quality matters more than speed.

Example: Product Launch Team



The key thing to remember is... Multi-agent systems are for complex work that benefits from specialized roles and built-in review cycles — like having an AI team instead of one AI assistant.

Choosing the Right Architecture

Use this **decision tree** to pick the simplest architecture that solves your problem — don't use a construction crew when you need a handyman.

The Decision Flowchart



APPENDIX

A. AI Tools Directory

A comprehensive, curated list of the most popular and useful AI tools organized by category.

1. Conversational AI & Chat Assistants

Tool Name	What does it do?	Tool Link
ChatGPT	OpenAI's conversational assistant that can chat, code, write, summarize, and brainstorm across domains. Industry-leading Large Language Model(LLM).	Visit
Claude	Anthropic's AI model focused on safe, interpretable, and creative conversations with high contextual reasoning and long context windows.	Visit
Gemini	Google's AI model integrated into Search and Workspace. Provides conversational, multimodal, and contextual help.	Visit
Perplexity	AI search engine combining live web data and LLM reasoning to give factual, cited answers. Great for research.	Visit
Microsoft Copilot	Microsoft's AI assistant integrated across Windows, Edge, Office 365, and Bing for productivity enhancement.	Visit

2. Content Writing & Marketing

Tool Name	What does it do?	Tool Link
Writesonic	AI content generation platform for blogs, ads, and marketing copy. Boosts productivity for writers and marketers.	Visit
Jasper	AI marketing platform for creating blog posts, social media, ads, and long-form content with brand voice consistency.	Visit
Copy.ai	AI-powered copywriting tool for marketing content, product descriptions, email campaigns, and social media posts.	Visit
Notion AI	AI assistant built into Notion for writing, summarizing, brainstorming, and organizing knowledge bases.	Visit
Supergrow	AI marketing platform to help grow leads, optimize campaigns, and accelerate audience engagement.	Visit
Social Sonic	Helps create, schedule, and optimize social media content using AI-driven insights.	Visit

3. Code & Development

Tool Name	What does it do?	Tool Link
Claude Code	AI-powered command-line coding agent that can autonomously handle complex development tasks, refactor code, and execute multi-step workflows.	Visit
GitHub Copilot	AI pair programmer that suggests code completions, entire functions, and helps debug. Powered by OpenAI Codex.	Visit
Cursor	AI-powered code editor built on VSCode with advanced code generation, refactoring, and chat capabilities.	Visit

Replit	Collaborative online IDE with AI code assistance for real-time coding and learning. Great for beginners.	Visit
Bolt	Developer tool or automation assistant built for fast prototyping and deployment of apps or workflows.	Visit
Tabnine	AI code completion tool supporting multiple languages and IDEs with privacy-focused options.	Visit
Codeium	Free AI-powered code acceleration toolkit with autocomplete, chat, and search across your codebase.	Visit

4. Design & Visual Content

Tool Name	What does it do?	Tool Link
Midjourney	Text-to-image model producing high-quality, artistic visuals for creators and designers. Industry-leading quality.	Visit
DALL-E 3	OpenAI's text-to-image generator integrated into ChatGPT, creating precise, contextual images from descriptions.	Visit
Stable Diffusion	Open-source text-to-image model allowing local deployment and customization for image generation.	Visit
Adobe Firefly	Adobe's AI image generator integrated into Creative Cloud for commercial-safe generative AI content.	Visit
Canva AI	AI-powered design features in Canva including Magic Design, background removal, and content generation.	Visit
Leonardo	AI art platform for creating game assets, illustrations, and concept art using text prompts.	Visit
Krea	AI design and art creation tool enabling rapid visual exploration and creative experimentation.	Visit
Magnific AI	AI image upscaler and enhancer that adds detail, improves resolution, and refines visuals.	Visit
Phot AI	AI-powered tool for editing, enhancing, and generating photos or visual content. Great for quick creative visuals.	Visit

5. Video Generation & Editing

Tool Name	What does it do?	Tool Link
Runway ML	Creative AI suite for video editing, image generation, and media production using machine learning.	Visit
HeyGen	AI video generator that turns text or scripts into realistic avatar videos with voice and lip sync.	Visit
Kling	Emerging generative video platform focusing on ultra-realistic, cinematic outputs.	Visit
Pika	AI video generation platform for creating and editing videos from text prompts with cinematic quality.	Visit
	AI-powered video and podcast editor with transcription, overdub, and multi-track	

Descript	editing capabilities.	Visit
Higgsfield	Advanced AI company developing realistic 3D / video generation technology for creative industries.	Visit

6. Audio & Voice

Tool Name	What does it do?	Tool Link
Eleven Labs	Industry-leading AI voice synthesis platform for lifelike text-to-speech and dubbing in multiple languages.	Visit
Suno	AI music generator for composing songs, jingles, and soundscapes from text prompts.	Visit
Whispr Flow	Converts speech to text in real time across apps. Helps users dictate, format, and edit content hands-free.	Visit
Murf AI	AI voice generator for creating professional voiceovers for videos, presentations, and e-learning.	Visit
Vapi	Voice or visual API platform enabling AI calling agents or multimodal experiences.	Visit

7. Research & Knowledge Management

Tool Name	What does it do?	Tool Link
Notebook LM	Google's AI research assistant that summarizes, queries, and connects your notes and documents intelligently.	Visit
Elicit	AI research assistant that helps find, summarize, and extract data from academic papers.	Visit
Consensus	AI-powered search engine that finds answers from scientific research with citations.	Visit
Semantic Scholar	AI-powered academic search engine helping researchers find relevant papers with intelligent filtering.	Visit
Chronicle	AI-powered tool for journaling, storytelling, or knowledge management to capture key moments.	Visit

8. Productivity & Automation

Tool Name	What does it do?	Tool Link
N8N	Open-source workflow automation tool with 400+ integrations, AI nodes, and self-hosting options for building complex automation pipelines.	Visit
Zapier AI	Workflow automation platform with AI-powered app integration and intelligent workflow suggestions.	Visit
Make (Integromat)	Visual automation platform with AI modules for connecting apps and automating complex workflows.	Visit
Mem	AI-powered note-taking and knowledge management that auto-organizes and surfaces relevant information.	Visit
Motion	AI calendar and project management tool that automatically schedules tasks and	Visit

	optimizes your day.	
Rocket	Automation tool that accelerates tasks, launches workflows, or optimizes processes using AI.	Visit
Numerous AI	A multi-purpose AI platform offering various automation and generation tools under one suite.	Visit

9. Meeting & Collaboration

Tool Name	What does it do?	Tool Link
Fireflies	Records, transcribes, and summarizes meetings automatically. Integrates with Zoom, Meet, and Teams to extract insights.	Visit
Otter.ai	AI meeting assistant that transcribes conversations in real-time, generates summaries, and extracts action items.	Visit
Granola	AI note-taking assistant for meetings — transcribes, summarizes, and organizes discussions.	Visit
Fathom	Free AI meeting assistant that records, transcribes, and summarizes video calls with instant highlights.	Visit

10. Data & Analytics

Tool Name	What does it do?	Tool Link
Lyzr AI	AI platform for analytics and automation — "laser-focused" insight generation and workflow optimization.	Visit
Julius AI	AI data analyst that helps analyze, visualize, and interpret data through natural language conversations.	Visit
DataChat	Conversational AI for data analytics, allowing teams to analyze data using natural language.	Visit

11. Translation & Language

Tool Name	What does it do?	Tool Link
DeepL	AI-powered translation service known for natural, accurate translations across 30+ languages.	Visit
Reverso	AI translation tool with context examples, grammar checking, and pronunciation features.	Visit

12. Specialized & Emerging Tools

Tool Name	What does it do?	Tool Link
Lovable	AI design assistant helping teams quickly create delightful, user-friendly web apps.	Visit
Emergent	AI discovery engine identifying emerging trends, ideas, and insights from large datasets.	Visit

Happenstance	AI idea generator fostering serendipitous discoveries, creative prompts, and connections.	Visit
Crystal	AI tool that analyzes personality and communication style to improve interpersonal effectiveness.	Visit
Humanic AI	Focuses on human-centric AI for personalization, empathy modeling, and user understanding.	Visit
Genspark	Generates creative ideas, articles, and media using generative AI — a "spark" for inspiration.	Visit
Emily	AI tool for engineers to scaffold, deploy, and manage ML or microservice projects. Simplifies orchestration and deployment.	Visit

B. Quick References / Cheat Sheets

A consolidated collection of all essential quick references for easy lookup.

B.1 How GenAI Works

Component	One-Liner
Training	Learn patterns from data
Parameters	Store patterns as numbers
Tokens	Break text into chunks
Attention	Decide what's relevant
Prediction	Guess next token, repeat
Temperature	Safe ↔ Creative dial

Key takeaway: GenAI doesn't think — it asks "What word probably comes next?" thousands of times. That's the whole trick.

B.2 AI Model Types

Term	One-Liner
Foundational	The base AI "engine" trained on massive data
Proprietary	Secret recipe — access via API only
Open Source	Public recipe — download and modify freely
Application	Products built ON TOP of foundational models
Perplexity	Search app using GPT-4/Claude underneath
Gamma	Presentation app using GPT-4/Claude underneath

Key takeaway: Foundational models are the engines (GPT-4, Llama). Proprietary = secret engine, Open source = public engine. Apps like Perplexity and Gamma are car service built using those engines — they add a nice interface but don't build the AI themselves.

B.3 Prompt Engineering Techniques

Technique	Complexity	Token Cost	Best For	Avoid For
Zero-Shot	Low	Low	Simple tasks	Complex formatting
Few-Shot	Medium	Medium	Pattern matching	Self-explanatory tasks
CoT	Medium	High	Multi-step reasoning	Simple lookups
ReAct	High	High	Agent workflows	Linear tasks
Role-Based	Low	Low	Domain expertise	Generic tasks
Chaining	High	High	RAG pipelines	Single-step tasks
Constraints	Low	Low	Compliance	Creative tasks
Self-Consistency	High	Very High	High-stakes decisions	Deterministic tasks
Meta-Prompting	Medium	Medium	Template building	Direct execution
Negative	Low	Low	Behavior prevention	Positive alternatives exist
Output Structuring	Medium	Low	API integration	Human-only reading
Iterative	High	High	Quality refinement	One-shot needs

B.4 AI Architectures

If Your Task Is...	Use This	Example
General question, no custom data	Basic LLM	"Explain blockchain"
Q&A over your documents	RAG	"What does our policy say about..."
Same steps every time	Workflow	Weekly report automation
Dynamic, figure-it-out task	Agent	"Research and summarize topic X"
Complex project needing review	Agentic	Content creation with editing cycles

Key takeaway: Start with the simplest option that works. You can always upgrade to a more complex architecture if needed.

B.5 General Concepts

Concept	One-Liner
Generative AI	Predicts what should come next based on patterns
Hallucinations	AI confidently making things up
Knowledge Cutoff	AI's "last updated" date
Context Window	How much AI can "remember" in a conversation
VERIFY	6-step checklist for responsible AI use
COSTAR	6-part prompting framework (Context, Objective, Style, Tone, Audience, Response)
Few-Shot	Teaching AI by showing examples

RAG	AI + your documents
Workflow	Fixed sequence of AI steps
Agent	AI that uses tools and decides next steps
Agentic	Team of AI agents working together

Remember: AI is a powerful tool, not magic. You're the pilot — AI is the autopilot. Know when to take the controls.